

# Developing Grammar Fuzzers

Aarush Kumbhakern, Ashoka University, India

## Abstract

Fuzzing is used for testing the resiliency of programs to input-based attacks. Fuzzers prioritize two key aspects - effectiveness in identifying vulnerabilities and the efficiency or speed at which they operate. While random fuzzers are exceptionally fast, their efficiency is hindered by input parsing challenges. In contrast, grammar fuzzers achieve deeper code penetration by generating specification conforming output.

This paper describes the development and optimization of a suite of grammar fuzzers. The paper includes comprehensive exploration of several fuzzing techniques, emphasizing efficiency through performance testing. Additionally, the paper delves into specific insights relevant to C language development and explores potential novel approaches to enhance existing fuzzer sub-processes.

## 1 Introduction

Fuzzing is a method used to test the resiliency of programs against unexpected, malformed, or malicious user input, by generating and injecting large amounts/iterations of such input. An example of a simple fuzzer is a random string generator. However, the majority of programs that take user input have input parsers to validate the input against a program-specific input language (e.g. a CSS minifier likely first checks if the input text is in fact valid CSS). In most instances, a randomly generated string would fail at this initial validation stage, and would be largely ineffective in testing deeper program logic. To reach program internals beyond the parser, a fuzzer must generate inputs that follow the grammar of the program's input language - such a fuzzer is termed a generation-based or grammar fuzzer [Felderer et al., 2016].

### 1.1 Grammars

A grammar consists of definitions (e.g. "<start>": [...]), which in turn consist of rules (e.g. ["<expr>"]), which then in turn consist of symbols/tokens (e.g. "<expr>") [Salem and Song, 2019]. A grammar contains two types of symbols:

1. **Terminal:** single ASCII characters, that when encountered are directly written to the fuzzer's output. Termed as such due to the fact that they "terminate" the expansion
2. **Nonterminal:** demarcated by surrounding angled brackets (<...>), each corresponding to a definition. Nonterminals are heuristically expanded into any of the corresponding definition's rules

The <start> token (or some alias) must exist in any grammar, and is the starting point of the expansion of a grammar. Listing 1 shows a JSON grammar that defines an arithmetic expression.

Listing 1: Arithmetic expression grammar

```
1 {
2   "<start>": [
3     ["<expr>"]
```

```
4   ],
5   "<expr>": [
6     ["<term>", "+", "<expr>"],
7     ["<term>", "-", "<expr>"],
8     ["<term>"]
9   ],
10  "<term>": [
11    ["<factor>", "*", "<term>"],
12    ["<factor>", "/", "<term>"],
13    ["<factor>"]
14  ],
15  "<factor>": [
16    ["+", "<factor>"],
17    ["-", "<factor>"],
18    ["<integer>", ".", "<integer>"],
19    ["<integer>"]
20  ],
21  "<integer>": [
22    ["<digit>", "<integer>"],
23    ["<digit>"]
24  ],
25  "<digit>": [
26    ["0"], ["1"], ["2"],
27    ["3"], ["4"], ["5"],
28    ["6"], ["7"], ["8"],
29    ["9"]
30  ]
31 }
```

### 1.2 Expansion

A grammar fuzzer, at its core, automates the expansion of a given grammar by stochastically expanding nonterminal tokens based on their corresponding definitions. Listing 2 shows pseudocode for a simple grammar fuzzer.

Listing 2: Simple grammar fuzzer pseudocode

```
1 def fuzz(grammar):
2   fuzz = "<start>"
3
4   while not fully_expanded(fuzz):
5     for token in fuzz:
6       if nonterminal(token):
7         expansion = expand(token, grammar)
8         fuzz.replace(token, expansion)
9
10  return fuzz
```

Listing 3 describes an example expansion trace of the arithmetic grammar, where each step expands all nonterminal tokens in the preceding step.

Listing 3: Arithmetic grammar expansion

```
1 <start>
2 <expr>
3 <term>
4 <factor>
5 <integer>.<integer>
6 <digit>.<digit><integer>
7 8.1<digit>
8 8.12
```

A notable issue with the simplistic approach in Listing 2 is the lack of control over expansion depth. Many grammars feature recursive rules, and without proper control, a simple grammar fuzzer may become ensnared in an infinite expansion loop. Moreover, even reaching a sufficiently large expansion depth can have

adverse effects on a recursive fuzzer, especially when constrained by a fixed stack size.

Beyond this challenge, the primary concern in simple fuzzing lies in speed (measured in terms of throughput in Kb/s, or output size over time), which is essential as a fuzzer is expected to rapidly generate either single inputs of considerable size or a substantial number of inputs of varying sizes.

### 1.3 Outline

After briefly describing evaluation standards and methods and providing hardware specifications of the test machine (Section 2), this paper is organized as follows:

**Section 3** We briefly discuss recursive and iterative approaches to how a fuzzer might handle expansion of a grammar.

**Section 4** We discuss various methods to control expansion depth specific to each control flow approach, with a specific focus on creating a mechanism to maintain a “current depth” state.

**Section 5** We create iterative and recursive variants of grammar-agnostic generic fuzzers, and discuss the construction of grammar structures for use with generic fuzzers.

**Section 6** We create Python compilers to compile an input grammar into C program code, condensing the generic fuzzers’ runtime grammar evaluation into a single compilation step to achieve significant increases in efficiency. We also discuss for the first time methods to algorithmically break recursive rings to generate cheap grammars.

**Section 7** We benchmark each fuzzer variant described in this paper according to a set of standards, and compare and discuss differences in performance statistics

## 2 Evaluation and Hardware

All fuzzers are evaluated on the basis of throughput in Kb/s.

All fuzzers achieve seeded output parity, ie. for a given seed value, all fuzzer variants will generate the same fuzzed string. Fuzzers will therefore be evaluated on their average performance at various depths for a fixed set of 10 randomly generated seeds.

The development and testing machine is a Zephyrus G14 (2021) running Fedora Linux on an AMD Ryzen 5800HS 8-core processor with a 2.8 GHz frequency, 512kb L2 cache, and 16MB shared L3 cache.

## 3 Expansion Control Flow

This section explores recursive and iterative approaches to the core fuzzing process of fully expanding a grammar - going from `<start>` to a fuzzed output string.

### 3.1 Recursive Expansion

Figure 1 represents lines 4 to 8 of the example expansion in Listing 3 as a tree.




Figure 1: Recursive expansion tree

Consider each nonterminal node to be a call to a `fuzz()` function (e.g. “Digit” is `fuzz(digit)`), and each terminal leaf to be the action of `print()`-ing that terminal token to the fuzzer’s output. Here, `fuzz()` writes terminals to the output with and generates recursive calls to `fuzz()` for each nonterminal token in the selected expansion. Listing 4 shows the static (pseudocode) function body of the `Factor` node in Figure 1 (note that this function body is specific to this particular call to `fuzz()`).

Listing 4: Recursive expansion pseudocode fragment

```

1 fuzz(factor):
2   fuzz(integer)
3   print(".")
4   fuzz(integer)
  
```

### 3.2 Iterative Expansion

Consider the definition of `expr` in the arithmetic grammar in 1. Two rules in this definition are self-referential, implying a recursive call to the fuzzer as seen in the previous section. Starting from `expr`, if each subsequent `expr` token in the expansion is expanded into a self-referencing rule, this will eventually lead to stack exhaustion as a result of a large number of recursive calls.

To address specifically the concern of a limited stack size, an alternative iterative approach can be adopted, wherein the program explicitly manages its own stack. In our pure iterative expansion method, the fuzzer iteratively mutates a *production string* or *expansion stack*. Table 1 shows the iterative version of the same detailed expansion described in Figure 1.

Stack	Out
<code>&lt;factor&gt;</code>	
<code>&lt;integer&gt;.&lt;integer&gt;</code>	
<code>&lt;digit&gt;.&lt;integer&gt;</code>	
<code>8.&lt;integer&gt;</code>	
<code>8.&lt;integer&gt;</code>	8
<code>&lt;integer&gt;</code>	8.
<code>&lt;digit&gt;&lt;integer&gt;</code>	8.
<code>1&lt;integer&gt;</code>	8.
<code>&lt;integer&gt;</code>	8.1
<code>&lt;digit&gt;</code>	8.1
<code>2</code>	8.1
	8.12

Table 1: Iterative Expansion

Observe that the number of steps in the expansion (excluding the pre-iteration first step) corresponds to the number of nodes in the recursive expansion tree (Figure 1).

It is essential to note that the fuzzer exclusively operates on the head of the expansion stack, writing to the fuzzer's output if the token is a terminal and substituting it with an expansion if the token is a nonterminal. Listing 5 shows pseudocode for the iterative expansion approach.

Listing 5: Iterative expansion pseudocode

```

1 def fuzz(stack):
2     while length(stack) > 0:
3         head, remainder = stack
4
5         if terminal(head):
6             print(head)
7             stack = remainder
8
9         else:
10            stack = expand(head) + remainder

```

### 3.3 A Common Approach

After examining the explicit expansion steps outlined in the iterative expansion (Table 1) and comparing them to the recursive tree representation of the same expansion (Figure 1), we observe that the expansion process follows an approach analogous to a depth-first left-to-right scan of the expansion tree. Notably, this approach is consistent across all fuzzers presented in this paper.

## 4 Limiting Expansion Depth

Referring again to the tree in Figure 1, the term *expansion path* will hereafter refer to a series of contiguous nodes (e.g. `factor`  $\rightarrow$  `integer`  $\rightarrow$  `digit`  $\rightarrow$  `"1"`) with reference to a single origin node.

Depth also is visually represented by the vertical position of tokens in the tree. If we consider `factor` to start at depth 1, both `digit` tokens, for example, are at a depth of 3 in their expansion paths.

### 4.1 Approach Overview

It is more accurate to term expansion depth as *free* expansion depth, ie. the expansion depth up to which the fuzzer is entirely uninhibited in its expansion of nonterminal tokens. Once an expansion path grows to a certain maximum expansion depth  $\mu$ , the fuzzer adopts an approach of selecting exclusively non-recursive (“cheap”) rules to continue the path. This renders any token  $x$  s.t.  $\text{depth}(x) = \mu$  functionally terminal.

To illustrate, consider  $\mu = 1$  and Figure 2, which frames a subtree of the expansion tree in Figure 1.

Assuming `factor` is at a depth of 0, the first `integer` in the subtree is observed at a depth of 1 ( $\geq \mu$ ) in the expansion path. Consequently, a depth-limited grammar fuzzer would exclusively choose “cheap” rules from the definition of `integer` to continue the expansion. The annotated definition of `integer` from the arithmetic grammar (Listing 1) in Listing 6 highlights the cheap and expensive rules.

Listing 6: Annotated integer Definition

```

1 "<integer>": [
2     ["<digit>", "<integer>"], # (expensive)
3     ["<digit>"]               # (cheap)
4 ],

```




Figure 2: Depth-limiting candidate

By assigning appropriate costs to each rule, only one possible expansion of `integer` remains, accordingly simplifying the framed subtree in Figure 2 into the subtree in Figure 3.




Figure 3: Limited subtree

Thus, to successfully limit expansion depth, two key requirements must be addressed:

1. Maintaining a “current depth” state specific to the fuzzer’s position in the current expansion pathway
2. Separately identifying cheap and expensive rules

At present we will assume manual identification and separation of rules by cost, and consider only direct recursion, ie. self-referencing rules. The paper will further explore this aspect of depth-limiting in Section 6.2, where we will algorithmically determine whether a rule is cheap or expensive, and additionally identify indirect recursion in grammars.

This section will focus solely on maintaining the “current depth” state of a fuzzer.

### 4.2 Recursive Limiting

The fact that a recursive approach explicitly constructs a tree makes maintaining a “current depth” state trivial. The process becomes straightforward as we only need to increment an argument for subsequent function calls, ie. `fuzz(x, depth)  $\rightarrow$  fuzz(y, depth + 1)`. Figure 4 demonstrates by modifying the tree in Figure 1.

Listing 7 shows the modified function body of the `Factor` node.

Listing 7: Recursive depth parameter

```

1 fuzz(factor, depth):
2     fuzz(integer, depth + 1)
3     print(".")
4     fuzz(integer, depth + 1)

```

### 4.3 Iterative Limiting

Maintaining a depth state for a pure iterative expansion approach is considerably more complex due to the implicit nature of the

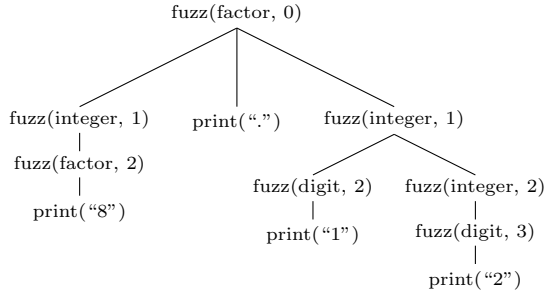


Figure 4: Recursive expansion tree

expansion, and the fact that the fuzzer just runs a loop which in itself provides no useful depth state.

Our solution uses a combination of a “history” array `stc` (Step-wise Token Count), and a (current) `depth` state variable. The history array acts as a step-wise token count, with each index corresponding to a specific depth in the expansion path. Simultaneously, the current depth state variable dynamically tracks the depth during the iterative process.

#### 4.3.1 Incrementing Depth

Incrementing depth is trivial - every time the iterative fuzzer encounters a nonterminal, we increment `depth`. We will begin constructing a stripped-down implementation of only the limiting portion of the code in C:

Listing 8: Rudimentary depth increment

```

1 while (STACK_LEN > 0) {
2   if(token >= 0) { // if terminal
3     continue;
4   }
5
6   depth++;
7 }

```

#### 4.3.2 Maintaining Token Count History

The key to figuring out when to decrement depth, is knowing when the expansion of a token in the current expansion path at a given depth is resolved, ie. all tokens in the expansion have been evaluated. This critical task is facilitated by the history array (`stc`).

Every time the fuzzer encounters a nonterminal, before prepending the nonterminal’s expansion to the production string, the token count of the expansion is written to the history array, and the current depth is then incremented as discussed above:

Listing 9: Writing token count to memory

```

1 while (STACK_LEN > 0) {
2   if(token >= 0) { // if terminal
3     continue;
4   }
5
6   // record token count then increment depth
7   stc[depth++] = rule->token_count;
8 }

```

#### 4.3.3 Evaluating Tokens

Every iteration of the fuzzer corresponds to the evaluation of a single token, underscoring the granularity of the iterative expansion process. In this context, with each iteration, the number of un-evaluated tokens in the expansion of the token at `depth - 1` is systematically decremented:

Listing 10: Every loop is an evaluation

```

1 while (STACK_LEN > 0) {
2   // decrement latest token count
3   if (depth > 0) stc[depth - 1]--;
4
5   if(token >= 0) { // if terminal
6     continue;
7   }
8
9   // record token count then increment depth
10  stc[depth++] = rule->token_count;
11 }

```

#### 4.3.4 Resolving Expansions

We recognize that a token is fully resolved only when the final terminal in its expansion is encountered. To illustrate, consider Figure 5:




Figure 5: Partial resolution example

In this example, `foo` is partially resolved when the fuzzer moves on to `bar` as all the tokens in the current expansion of `foo` have been evaluated. However, `foo` achieves complete resolution only upon encountering the final terminal of its full expansion “c”.

A fundamental understanding emerges: a nonterminal will always increase the depth, and only a terminal can signal a decrease in depth. Even if a token becomes fully resolved by the evaluation of token  $x$ , if  $x$  is a nonterminal, decreasing the depth would be premature. We therefore check for resolution only upon encountering a terminal. Listing 11 demonstrates the resolution check:

Listing 11: Checking for resolution

```

1 while (STACK_LEN > 0) {
2   // decrement latest token count
3   if (depth > 0) stc[depth - 1]--;
4
5   if(token >= 0) { // if terminal
6     // rollback to nearest unresolved
7     while (stc[depth - 1] == 0) depth--;
8
9     continue;
10  }
11
12  // record token count then increment depth
13  stc[depth++] = rule->token_count;
14 }

```

In this code snippet, the check for resolution is placed in a conditional `while` loop. This design ensures that complete resolution

cascades down through any preceding partial resolutions. For instance, in the expansion of `foo`, encountering `"c"` marks not only `bar` as resolved but cascades down through the partial resolution of `foo`, marking it as completely resolved as well.

#### 4.3.5 Accounting for Limiting

As we introduce limiting, the logic of the depth state mechanism undergoes a slight modification. Recall that any token  $x$  at the maximum depth  $\mu$  becomes functionally terminal due to the limiting. Consequently, at this point, the depth state mechanism focuses solely on whether  $x$  has been fully resolved. Any expansions or sub-expansions of  $x$  are flattened, and depth updates are temporarily paused.

This adjustment in the depth state mechanism reflects the shift in focus when limiting is introduced. The mechanism now prioritizes the determination of full resolution for tokens at the maximum depth, simplifying the expansion process during this specific phase.

Listing 12: Complete depth state maintenance mechanism

```
1 while (STACK_LEN > 0) {
2     // decrement latest token count
3     if (depth > 0 && depth < max_depth) stc[depth - 1]--;
4
5     if(token >= 0) { // if terminal
6         // if haven't already, decrement latest token count
7         if (depth > 0 && depth >= max_depth) stc[depth - 1]--;
8
9         // rollback to nearest unresolved
10        while (stc[depth - 1] == 0) depth--;
11
12        continue;
13    }
14
15    // record token count then increment depth
16    if (depth < max_depth) stc[depth++] = rule->token_count;
17    else stc[depth - 1] += (rule->token_count) - 1;
18 }
```

#### 4.3.6 Example Execution

Table 2 shows an example of the iterative depth state mechanism in action, using the iterative expansion example in Section 3.2.

Depth	stc[]	Stack	Out
0	[]	<factor>	
1	[3]	<integer>.<integer>	
2	[2,1]	<digit>.<integer>	
3	[2,0,1]	8.<integer>	
1	[2,0,0] → [2]	.<integer>	8
1	[1]	<integer>	8.
2	[0,2]	<digit>.<integer>	8.
3	[0,1,1]	1.<integer>	8.
2	[0,1,0] → [0,1]	<integer>	8.1
3	[0,0,1]	<digit>	8.1
4	[0,0,0,1]	2	8.1
0	[0,0,0,0] → []		8.12

Table 2: Iterative Depth State Mechanism

## 5 Generic Fuzzers

A generic fuzzer is a function that takes a grammar defined in a certain predefined structure as an input, and produces a fuzzed output that conforms to said grammar. Evaluation and expansion of the grammar are both runtime processes.

In this section we construct recursive and iterative variants of a generic fuzzer in C. We will begin by defining the structure for the input grammar.

### 5.1 Grammar Structure

We first declare the components that will comprise the grammar - tokens, rules, and definitions - and a simple structure for the grammar itself:

Listing 13: Generic grammar components

```
1 typedef signed char token_t;
2
3 typedef struct Rule {
4     size_t token_count;
5     token_t* tokens;
6 } Rule;
7
8 typedef struct Definition {
9     size_t rule_count[2];
10    Rule* rules[2]; // [0]: cheap, [1]: costly
11 } Definition;
12
13 typedef struct Grammar {
14     Definition* definitions;
15 } Grammar;
```

The struct components are self-explanatory, but it is worth clarifying with regard to the rules member in Definition that `rule[0]` contains the definition's cheap rules, and `rule[1]` contains its expensive rules.

The typedef for tokens is specifically of type `signed char` as we will be representing terminals with their corresponding (positive) ASCII codes, and nonterminals with negative values - this allows identifying nonterminals to be simplified to a trivial arithmetic inequality test. The way these negative values are assigned is as follows (using the arithmetic grammar (Listing 1) as an example):

Listing 14: Defining nonterminals

```
1 enum nonterminals {start = SCHAR_MIN, expr, term, factor,
2     integer, digit};
```

The `start` token is assigned the lowest possible negative value, and each key in the grammar is automatically assigned the next lowest value in order of appearance in the `enum`. The value of each key is then tied to the index of its corresponding definition in the final grammar using designated initializers.

Listing 15: Generic grammar snippet

```
1 Grammar grammar = {.definitions=(Definition []) {
2     [start - start] = (Definition) { // start
3         .rule_count={1, 0},
4         .rules={
5             (Rule []) { // "<rule>"
6                 (Rule) {
7                     .token_count=1,
8                     .tokens=(token_t[]) {expr}
```

```

9         },
10        },
11    }
12 },
13
14 [expr - start] = (Definition) { // expr
15     .rule_count={1, 2},
16     .rules={
17         (Rule []) {
18             (Rule) { // "<term>"
19                 .token_count=1,
20                 .tokens=(token_t[]) {term}
21             },
22         },
23         (Rule []) {
24             (Rule) { // "<term>+<expr>"
25                 .token_count=3,
26                 .tokens=(token_t[]) {term, '+', expr}
27             },
28             (Rule) { // "<term>-<expr>"
29                 .token_count=3,
30                 .tokens=(token_t[]) {term, '-', expr}
31             },
32         },
33     }
34 },
35
36 // ...and so on
37 } };

```

## 5.2 Recursive Generic Fuzzer

### 5.2.1 Making Recursive Calls

In the recursive expansion trees we have seen so far, tokens are directly expanded via recursive calls into a chosen expansion, i.e. for  $\langle x \rangle \rightarrow \langle a \rangle + \langle b \rangle$ , we would generate the pseudocode in Listing 16.

Listing 16: Recursive expansion

```

1 fuzz(x):
2     fuzz(a)
3     print("+")
4     fuzz(b)

```

However, actual implementation requires the fuzzer to first select a rule and then, depending on the type of token, make recursive sub-calls or write to output for each nonterminal in the expansion, all within a loop, as shown in Listing 17.

Listing 17: Realistic evaluation

```

1 fuzz(x):
2     rule = get_rule(x) // [a, "+", b]
3
4     for token in rule:
5         if rule is terminal: print(token)
6         else: fuzz(token)

```

Listing 18 shows the logic up to this point translated into C.

Listing 18: Rudimentary implementation

```

1 fuzzer(Grammar* grammar, token_t token) {
2     Definition* def = &grammar->definitions[token - start];
3     Rule* rule = get_rule(def)
4
5     for (size_t i = 0; i < rule->token_count; i++) {
6         // if terminal write to stdout

```

```

7         if (rule->tokens[i] >= 0) {
8             putchar(rule->tokens[i]);
9         } else { // otherwise make recursive call
10             fuzzer(grammar, rule->tokens[i]);
11         }
12     }
13 }

```

We will discuss the actual implementation of `get_rule()` in section 5.2.3 below - for now, we simply assume that it returns a random rule from the definition passed to it.

### 5.2.2 Limiting Depth

Limiting depth for a recursive, as discussed, is primarily an issue of rule selection and depth state maintenance.

We will assume that `get_rule(token, cheap)` will choose only cheap rules if `cheap = 1`.

The condition for switching to an exclusively cheap rule selection mode is `depth >= max_depth`. `max_depth` can be passed in as a static parameter, and as discussed in Section 4.2 depth may be simply passed as an argument for a recursive fuzzer. Listing 19 shows the updated recursive fuzzer.

Listing 19: Depth-limited generic recursive fuzzer

```

1 fuzzer(Grammar* grammar, depth_t max_depth, depth_t depth,
2         token_t token) {
3     Definition* def = &grammar->definitions[token - start];
4     Rule* rule = get_rule(def, depth >= max_depth)
5
6     for (size_t i = 0; i < rule->token_count; i++) {
7         // if terminal write to stdout
8         if (rule->tokens[i] >= 0) {
9             putchar(rule->tokens[i]);
10        } else { // otherwise make recursive call
11            fuzzer(grammar, max_depth, depth + 1,
12                  rule->tokens[i]);
13        }
14    }
15 }

```

An additional check has been added to ensure that the fuzzer does not attempt to choose expensive rules if the current token's definition does not contain any.

This completes the generic recursive fuzzer implementation.

### 5.2.3 Selecting Rules

The `get_rule()` function chooses a random rule from a given definition. If `cheap = 1`, it chooses only cheap rules.

Listing 20: Get rule function

```

1 Rule* get_rule(Definition* definition, int cheap) {
2     size_t cheap_c = definition->rule_count[0];
3     size_t exp_c = definition->rule_count[1];
4
5     // force cheap if no expensive rules
6     if (!exp_c) cheap = 1;
7
8     if (cheap) {
9         return &definition->rules[0][rand() % cheap_c];
10    } else {
11        // choose from cumulative rule set
12        size_t choice = rand() % (cheap_c + exp_c);

```

```

14     size_t cost = choice >= cheap_c;
15     if (cost) choice -= cheap_c;
16     return &definition->rules[cost][choice];
17 }
18 }

```

## 5.3 Iterative Generic Fuzzer

### 5.3.1 Iterative Expansion

To summarize the approach described in Section 3.2, a pure iterative fuzzer mutates an expansion stack by iteratively evaluating the first token in the stack - writing it to output if terminal and expanding it if nonterminal.

We will again begin by minimally implementing an iterative fuzzer that conforms to this description.

Listing 21: Simple generic iterative fuzzer

```

1 void fuzzer(Grammar* grammar) {
2     token_t* stack = malloc(1024 * sizeof(token_t));
3     stack[0] = start; // initialize stack with start
4
5     size_t stack_len = 1;
6
7     while (stack_len > 0) {
8         token_t token = stack[0]; // get first token
9
10        // if token is terminal write to stdout
11        if (token >= 0) {
12            // remove token from stack
13            memmove(stack, stack + 1, --stack_len);
14
15            putchar(token);
16
17            continue;
18        }
19
20        Definition* def = &grammar->definitions[token - start];
21        Rule* rule = get_rule(def);
22
23        // substitute token with rule in stack
24        memmove(stack + rule->token_count, stack + 1, stack_len
25                * sizeof(token_t));
26        memcpy(stack, rule->tokens, rule->token_count);
27        stack_len += rule->token_count - 1;
28    }
29 }

```

The iterative fuzzer shares the implementation of its `get_rule()` function with the recursive fuzzer (Section 5.2.3).

### 5.3.2 Limiting Depth

As it happens, we have already discussed everything needed to implement depth limiting in the iterative fuzzer. With the understanding that we are using the same `get_rule()` function, we can simply copy the internal rule choice code from the recursive fuzzer, and then merge the iterative depth limiting mechanism from Section 4.3 with the current code for the fuzzer:

Listing 22: Depth-limited generic iterative fuzzer

```

1 void fuzzer(Grammar* grammar, depth_t max_depth) {
2     token_t* stack = malloc(1024 * sizeof(token_t));
3     stack[0] = start; // initialize stack with start
4
5     size_t stack_len = 1;

```

```

6
7     // depth state variables
8     depth_t stepwise_token_count[128] = {};
9     depth_t current_depth = 0;
10
11    while (stack_len > 0) {
12        token_t token = stack[0]; // get first token
13
14        // decrement latest token count
15        if (depth > 0 && depth < max_depth) stc[depth - 1]--;
16
17        // if token is terminal write to stdout
18        if (token >= 0) {
19            // remove token from stack
20            memmove(stack, stack + 1, --stack_len);
21
22            putchar(token);
23
24            // if haven't already, decrement latest token count
25            if (depth > 0 && depth >= max_depth) stc[depth - 1]--;
26
27            // rollback to nearest unresolved token
28            while (stc[depth - 1] == 0) depth--;
29
30            continue;
31        }
32
33        Definition* def = &grammar->definitions[token - start];
34        Rule* rule = get_rule(def, depth >= max_depth);
35
36        // record token count then increment depth
37        if (depth < max_depth) stc[depth++] = rule->token_count;
38        else stc[depth - 1] += (rule->token_count) - 1;
39
40        // substitute token with rule in stack
41        memmove(stack + rule->token_count, stack + 1, stack_len
42                * sizeof(token_t));
43        memcpy(stack, rule->tokens, rule->token_count);
44        stack_len += rule->token_count - 1;
45    }

```

This completes the generic iterative fuzzer.

## 6 Compiled Fuzzers

To reiterate, generic fuzzers perform two major runtime processes:

1. Evaluate the input grammar
2. Expand the start token in a manner defined by the input grammar to create fuzzed output

This section aims to enhance fuzzer speed by consolidating the runtime evaluation of the input grammar into a single compilation step. The approach involves compiling a given grammar into program code specific to that grammar.

Since the compilation is a one-time activity and not a runtime process, speed is not a primary concern. For convenience, the compiler and its utilities will be implemented in Python but will output program code in C.

Moreover, this section will not delve into the details of the compiler's implementation but will instead concentrate on the expected output of the compiler. The compiler program will produce a C module, comprising a header file and an implementation file. However, our focus in this section will be solely on the im-

plementation file.

## 6.1 Generalized Pseudo-Implementation

### 6.1.1 Basic Compilation

The compilation process is straightforward. Definitions are transformed into functions, rules are selected by conditional matching with a randomly generated number within the range of the definition’s rule count, and each token in a rule is converted into a call to the corresponding definition if nonterminal, or a write function if terminal.

Consider the example in Listing 23 of the compilation of `expr` from the arithmetic grammar (Listing 1) into pseudocode.

Listing 23: Example compilation (pseudocode)

```

1  expr():
2    choice = random(3)
3
4    if choice == 0:
5      term()
6      print("+")
7      expr()
8
9    if choice == 1:
10     term()
11     print("-")
12     expr()
13
14    if choice == 2:
15     term()

```

### 6.1.2 Limiting Depth

A compiled fuzzer uses the same recursive approach to depth state maintenance as discussed in Section 4.2. Actual limiting via cheap rule section is achieved by compiling two variants - cheap and random - for each definition. When the fuzzer reaches a maximum depth, the random variant will redirect to the cheap variant.

Listing 24: Compiled fuzzer depth limiting

```

1  # cheap variant
2  expr_cheap():
3    choice = random(1)
4
5    if choice == 0:
6      term_cheap()
7
8
9  # random variant
10 expr_rand(max_depth, depth):
11   if depth >= max_depth: # redirect to cheap
12     expr_cheap()
13     return
14
15   choice = random(3)
16
17   if choice == 0:
18     term_rand(max_depth, depth + 1)
19     print("+")
20     expr_rand(max_depth, depth + 1)
21
22   if choice == 1:
23     term_rand(max_depth, depth + 1)
24     print("-")

```

```

25     expr_rand(max_depth, depth + 1)
26
27     if choice == 2:
28       term_rand(max_depth, depth + 1)

```

Tokens in the rules of the cheap variant of a generator only redirect in turn to the cheap generators of said tokens, ie. once the fuzzer switches to using a cheap generator, it exclusively uses cheap generators for the remainder of the expansion path.

The cheap variants are generated from a cheap sub-grammar of the original grammar. The next section will discuss how cheap grammars are generated.

## 6.2 Cheap Grammar Generation

A cheap grammar is a sub-grammar that contains only cheap (non-recursive) rules. The central operation in the cheapening of a grammar, and the focus of this section, lies in the identification and pruning of recursive rules. Up until now, the paper has assumed the manual identification and separation of cheap and expensive rules. However, it is imperative for us to develop the capability to algorithmically identify recursive rules in grammars.

### 6.2.1 Indirect Recursion

Up to this point, we have been exposed only to direct (self-referential) recursion. We will now broaden our definition of recursive rules to include indirect recursion. Consider the “perfidious” grammar in Listing 25.

Listing 25: Perfidious grammar

```

1  {
2    "<start>": [
3      ["<foo>"],
4      ["a"]
5    ],
6    "<foo>": [
7      ["<bar>"],
8      ["b"]
9    ],
10   "<bar>": [
11     ["<baz>"],
12     ["<start>"],
13   ],
14   "<baz>": [
15     ["(", "<foo>", ")"],
16     ["c"],
17   ]
18 }

```

The grammar contains no direct recursion, but further examination reveals two recursive “rings”:




Figure 6: Perfidious grammar rings

Our goal is to identify and break these rings, ie. ensure no such rings exist in the cheap grammar.



### 6.2.2 Cheapening approaches

There are two approaches to cheapening a grammar.

In a **permissive** approach, a rule is retained as long as the number of non-recursive expansion pathways is  $> 1$ , even all other expansion pathways are recursive. A permissive approach will only remove “vicious” rings, ie. rings that cannot be terminated without completing a full circuit. In other words, this approach will remove a rule if *all* expansion paths from said rule lead back to itself:

This means that rings may still exist in the grammar - for example, a permissive approach would not remove any of the rings in the recent perfidious grammar (Listing 25). This undermines the purpose of cheapening to the end of limiting depth, and does not conform to our “functionally terminal” (Section 4.1) requirement. Consequently, we proceed to explore the next approach.

In broad terms, an **aggressive** approach adopts a more stringent criterion by removing any rule that causes recursion. There are two sub-approaches to identifying the cause of a ring. To illustrate, consider Figure 7 - the expansion path diagram for the perfidious grammar (Listing 25).




Figure 7: Perfidious grammar expansion map

The first sub-approach would attribute recursion to the point of entry to a ring. In this **head-first** approach, we iterate over every rule in the grammar and prune it if we find any expansion paths leading from the rule that form a ring (i.e., the rule is an entry point).

Consider ring 1 in Figure 6, the entry points to which could be either  $\text{start} \rightarrow \text{foo}$  or  $\text{foo} \rightarrow \text{bar}$ . Figures 8 and 9 show the effects of pruning these entry points.




Figure 8: Pruning  $\text{start} \rightarrow \text{foo}$




Figure 9: Pruning  $\text{foo} \rightarrow \text{bar}$

As these examples illustrate, a head-first approach is prone to premature amputation of the grammar. Consequently, we discard

this approach in favor of the one recommended in this paper for cheapening a grammar.

In the paradigm of a **tail-first** approach, the true cause of a ring is not the entry point but rather the ‘re-entry’ point—the point at which the expansion path loops back onto itself. It is only at this point that the ring is formed.

Unlike the permissive and head-first approaches, where the operative entity is a rule, the operative entity of a tail-first approach is a pathway. Instead of saying “start is recursive”, we state “ $\text{bar} \rightarrow \text{start}$  makes  $\text{start} \rightarrow \text{foo} \rightarrow \text{bar} \rightarrow \text{start}$  recursive”. A tail first approach therefore requires a sequential evaluation of every expansion path from start.

Consider the simple example of a tail-first evaluation in action, in Figure 10. Using this approach, we flag blameful rules, and simply filter them out when creating a copy of the original grammar.

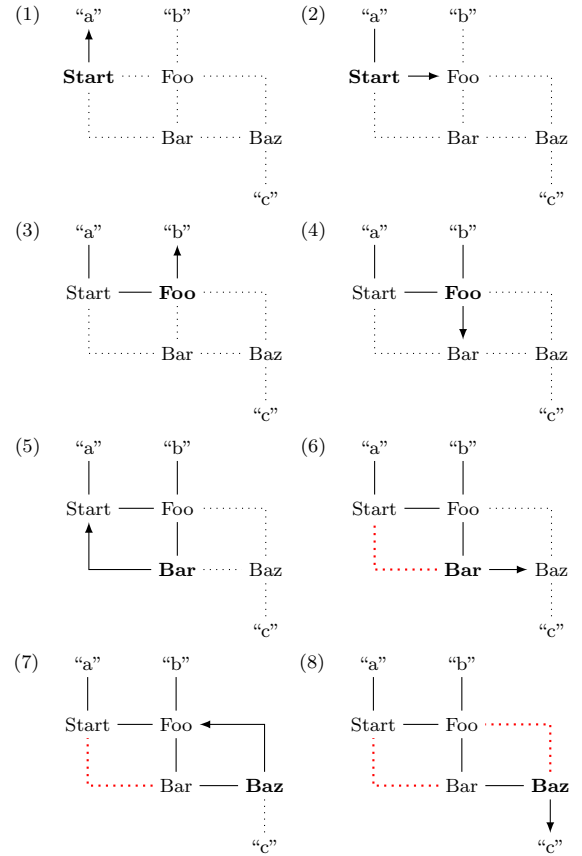


Figure 10: Tail-first evaluation of perfidious grammar

### 6.2.3 Cheapening Implementation

As discussed in the previous section:

A tail first approach [...] requires a sequential evaluation of every expansion path from start.

We will therefore begin by creating a function that traverses every expansion path from start. The approach is straightforward:

for a given nonterminal, obtain the list of rules in the corresponding definition, and for each rule, generate a recursive call to the traversal function for every token in the rule:

Listing 26: Grammar traversal

```
1 def flag(grammar, key="<start>"):
2     for rule in grammar[key]:
3         nonterms = set(nonterminals(rule))
4
5     for token in nonterms:
6         flag(grammar, token)
```

We also convert the list of nonterminals in a given rule into a set of unique values, to avoid redundant traversal.

To identify when an expansion path loops back onto itself, we must maintain a path-specific history. Since this is a recursive traversal, this is once again as trivial as inheriting a parameter:

Listing 27: Visited state

```
1 def flag(grammar, key="<start>", visited=["<start>"]):
2     for rule in grammar[key]:
3         nonterms = set(nonterminals(rule))
4
5     for tkn in nonterms:
6         flag(grammar, tkn, visited.copy() + [token])
```

Note the `copy()` method used when passing the visited to recursive sub-calls - this is to ensure the visited state is unique to each pathway, as Python by default passes lists by reference.

The final step is to actually identify a re-entry/loopback point and flag the offending rule. This is done simply by checking if any of the nonterminals in a given rule have previously been visited in the current expansion path:

Listing 28: Flagging re-entry points

```
1 def flag(grammar, key="<start>", visited=["<start>"],
2         flagged=[]):
3     for rule in grammar[key]:
4         if f"{key},{rule}" not in flagged:
5             nonterms = set(nonterminals(rule))
6
7             # if loopback, flag rule
8             if set(nonterms).intersection(set(visited)):
9                 flagged.append(f"{key},{rule}")
10
11            else: # otherwise continue traversal
12                for token in nonterms:
13                    flag(grammar, token, visited.copy() + [token],
14                        flagged)
15
16    return flagged
```

You will notice that `copy()` is not used when inheriting `flagged` - this is because the list of flagged rules is not specific to an expansion path, and so passing by reference (ie. sharing the list between all recursive calls to `flag()`) is desired behavior in this case.

To cheapen a grammar, `flag()` is used to identify recursive re-entry points, and the original grammar is simply duplicated leaving out any rules present in the output of `flag()`.

## 6.3 Recursive Compiled Fuzzer

We will now discuss the compiled output of a compiler that generates a recursive fuzzer.

### 6.3.1 Definitions to Generator Functions

Going from the pseudocode presented in Section 6.1 to actual C program code for a definition is fairly trivial:

Listing 29: Recursive compiled fuzzer snippet

```
1 // cheap variant
2 void expr_cheap() {
3     switch (rand() % 1) {
4         case 0:
5             term_cheap();
6             return;
7     }
8 }
9
10 // random variant
11 void expr_rand(int max_depth, int depth) {
12     if (depth >= max_depth) { // redirect to cheap
13         expr_cheap();
14         return;
15     }
16
17     switch (rand() % 3) {
18         case 0:
19             term_rand(max_depth, depth + 1);
20             return;
21
22         case 1:
23             term_rand(max_depth, depth + 1);
24             putchar(43);
25             expr_rand(max_depth, depth + 1);
26             return;
27
28         case 2:
29             term_rand(max_depth, depth + 1);
30             putchar(45);
31             expr_rand(max_depth, depth + 1);
32             return;
33     }
34 }
```

Each definition in a grammar will be compiled as shown above and written to a file (e.g. `core.c`).

### 6.3.2 Generating the Initiator

We will also need to generate and include a main `fuzz` function that will begin the fuzzing process - this will be common to any grammar:

Listing 30: Recursive compiled fuzzer initiator

```
1 void fuzz(int seed, int max_depth) {
2     // initialize random generator
3     srand((unsigned) seed);
4
5     gen_start_rand(max_depth, 0);
6 }
```

This completes the compiled recursive fuzzer.

## 6.4 Seeding (Interlude)

You will have noticed the introduction of a `seed` parameter to initialize the random generator in the previous section. This is not

exclusive to compiled fuzzers - the random generator of a generic fuzzer is initialized in the `main()` function where the `fuzzer()` function is called.

Due to the homologous approaches (Section 3.3) to fuzzing the fuzzer variants presented in this paper follow, passing the same seed to any variant will result in each fuzzer producing the same fuzzed output.

## 6.5 Iterative Compiled Fuzzer

The iterative fuzzer is a hybrid between the generic iterative fuzzer and the compiled recursive fuzzer in that while it maintains an expansion stack that it continually evaluates the head of, expansions are still sub-calls (albeit not recursive), allowing us to borrow the simple depth state maintenance mechanism from the recursive fuzzer.

Unlike the generic iterative fuzzer, however, the stack stores function calls complete with arguments, a design choice influenced by the compilation of definitions into generator functions outlined in previous sections.

### 6.5.1 Generating the Driver

We will begin this time by generating the driver function:

Listing 31: Iterative compiled fuzzer driver

```
1  Lambda stack[1024];
2  int stack_len = 1;
3
4  void fuzz(int seed, int max_depth) {
5      // initialize random generator
6      srand((unsigned) seed);
7
8      // initialize stack with call to start generator
9      stack[0] = (Lambda) {.args={max_depth, 0},
10         .func=&gen_start_rand};
11
12     // execute first call in stack until stack is empty
13     while (stack_len > 0) stack[0].func(stack[0].args[0],
14         stack[0].args[1]);
15
16     return;
17 }
```

`stack` and `stack_len` are declared with global scope so that they may be later accessed by generator functions.

### 6.5.2 Lambda Functions

You will notice that the stack is of type `Lambda`. This is a structure we have created for the purpose of storing function calls in the stack, comprised of a function pointer, and an argument array. Listing 32 shows the implementation of the structure, found in the header file generated by the compiler.

Listing 32: C Lambda Structure

```
1  typedef void (*func)();
2
3  typedef struct Lambda {
4      int args[2];
5      func func;
6  } Lambda;
```

Calling and creating a lambda are both demonstrated in the driver code in the previous section.

### 6.5.3 Definitions to Generator Functions

The compilation of definitions is very similar to that of the recursive compiled fuzzer, with one main modification - instead of directly calling the subsequent generator functions present in the chosen rule, we prepend them to the expansion stack similarly to the generic iterative fuzzer.

Consider Listing 33, showing this modification made to a rule from the recursive compiled fuzzer in Listing 29 (lines 23-26):

Listing 33: Iterative compiled fuzzer rule

```
1  // shift stack up to make space for expansion
2  memmove(&stack[3], &stack[1], stack_len * sizeof(Lambda));
3
4  // write function calls into space created
5  stack[0] = (Lambda) {.args={max_depth, depth + 1},
6     .func=&gen_term_rand};
7
8  stack[1] = (Lambda) {.args={43}, .func=&write};
9
10 stack[2] = (Lambda) {.args={max_depth, depth + 1},
11     .func=&gen_expr_rand};
12
13 stack_len += 2;
14 return;
```

### 6.5.4 Handling Nonterminals

The iterative compiled fuzzer, as you may have noticed in previous section's code, requires us to create a custom `write()` function to handle terminals. Where the recursive variant allowed us to simply use `putchar()`, additional stack operations in the iterative variant require a wrapper:

Listing 34: Custom write function

```
1  void write(int token) {
2      // remove call from stack
3      memmove(&stack[0], &stack[1], stack_len * sizeof(Lambda));
4
5      // write to stdout
6      putchar(token);
7
8      stack_len -= 1;
9      return;
10 }
```

This completes the compiled iterative fuzzer.

## 7 Profiling

Fuzzer variants of each class (generic/compiled) and subtype (recursive/iterative) were generated for the JSON, TinyC, and HTTP grammars. To profile each variant, the corresponding executable was run for seed values  $x$  such that  $\forall x \ 0 < x \leq 5000, x \in \mathbb{Z}$ . Each data point is the average throughput of 10 runs of the fuzzer for the corresponding seed value  $x$ . The max depth argument for each execution is set to 128.

## 7.1 Fuzzer Behavior

Figures 11 through 13 plot raw profiling data for the compiled recursive fuzzers. The purpose of these figures is to visualize the behavior of all fuzzer variants for a given grammar - specifically in terms of throughput and output size.

*Note: The behavior of each fuzzer class and subtype was observed to be identical for a given grammar, and as such we have chosen to include visualizations for only the compiled recursive variants.*




Figure 11: Compiled recursive JSON fuzzer - Size vs. Throughput




Figure 12: Compiled recursive TinyC fuzzer - Size vs. Throughput

Cursory analysis of the raw profiling data shows a drastically differing behavior profile for the TinyC grammar (Figure 12). We start by noting the possible correlation with average output




Figure 13: Compiled recursive HTTP fuzzer - Size vs. Throughput

size -  $2.44 \cdot 10^{-1} < 6.45 \cdot 10^2 \gg 5.89 \cdot 10^{-4}$  for JSON, TinyC, and HTTP respectively.

We then conjecture then that the TinyC plot (Figure 12) visualizes the “complete” behavior profile of the fuzzers in this paper, or conversely that the JSON and HTTP behavior profiles are simply a subset of the more expansive (on the basis of output size size) TinyC profile.

This is supported by Figure 14, which re-visualizes the behavior profile of the TinyC fuzzers by restricting output size to  $< 15$  kilobytes to approximate the output size range of the JSON fuzzer. We observe a distinct similarity between the truncated TinyC profile and the complete JSON profile - further limiting the visualization by an order of magnitude would likely yield a profile similar to the full HTTP profile.

In terms of the relationship between size and throughput, we also observe the following rough tendency:

$$\lim_{\text{size} \rightarrow 0} \text{throughput} \approx 0 \quad (1)$$

This suggests a lower limit on the execution time of the fuzzer that is independent of the output size, and is supported by the fact that each fuzzer variant profiled, regardless of grammar, had a minimum execution time  $x$  such that  $5 \cdot 10^{-4} < x < 6 \cdot 10^{-4}$ .

## 7.2 Benchmarking

In consideration of the correlation described in (1), for comparative benchmarking of fuzzer variants we considered the average throughput of the ten largest fuzzed outputs for each fuzzer. Recall that all fuzzers achieve seeded output parity, meaning this selection is consistent across all variants for a given grammar.

Figure 15 compares average maximal throughput for each fuzzer class/subtype, grouped by language.

We observe from the JSON and TinyC grammars that at large




Figure 14: Compiled recursive TinyC fuzzer - Size vs. Throughput (truncated)

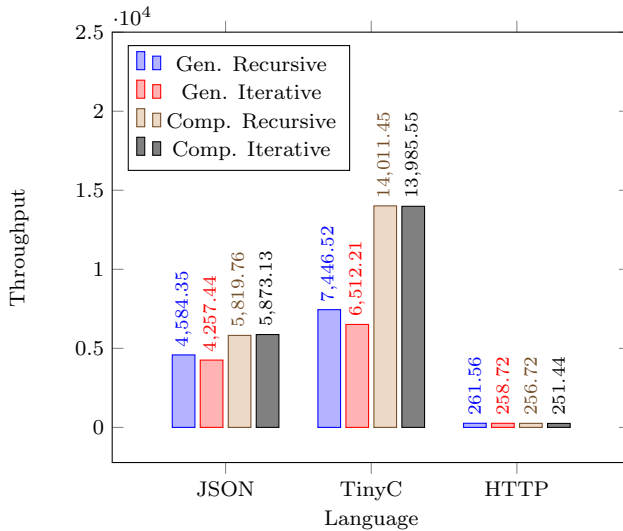


Figure 15: Maximal Throughput Comparison

“enough” average output sizes, compiled fuzzers consistently outperform the generic class variants. We also observe that the performance difference between subtypes of the same class is far more prominent in generic fuzzers than compiled.

The HTTP grammar, however, suggests that at low output sizes, all aforementioned differences become negligible (at least), with the generic variants slightly outperforming the compiled variants.

## 8 Discussion

We have now successfully developed generic and compiled fuzzers with both iterative and recursive flow control that can efficiently generate fuzzed output given a JSON grammar. Comparing the maximal throughput we achieved for the JSON grammar with

fuzzers described in this paper (5873.13 KB/s) with existing fast fuzzers for the same grammar<sup>1</sup> such as the F1 fuzzer [Gopinath and Zeller, 2019] (4281.5 Kb/s), Grammarinator [Hodován et al., 2018] (7 Kb/s), and an industry fuzzers such as Dharma by Mozilla [Diehl, 2014] (121.8 Kb/s), we find the performance of our fuzzers extremely adequate.

All in all, we have adequately demonstrated and designed fuzzers that are efficient in two ways - the first being in terms of raw throughput, and the second being in terms of *valid* throughput, by virtue of being grammar fuzzers. The fuzzers we’ve developed allow us to quickly generate large (or, as large as the grammar is designed to allow) inputs in high volumes, greatly optimizing the input generation process of fuzz testing.

## 9 Limitations and Future Work

Approaching the end of the scope of this paper, we now discuss points of note that extend outside it that have a bearing on the usefulness of our results.

### 9.1 Analysis

With what limited analysis we performed on the fuzzers, we have identified a key bottleneck in the form of minimum execution time that leads to the throughput of any fuzzer to follow a plateau curve against size. We have roughly conjectured that the full extent of the behavior each fuzzer described in this paper follows this curve, or a subset of it, depending on the size distribution of its output. Further analysis of this with a larger number of grammars and a more powerful testing machine would likely give us more detailed insight into the behavior of the fuzzers and the cause of the lower-limit bottleneck respectively.

### 9.2 Advanced Features

The fuzzers developed in this paper are on the lower end of complexity as far as grammar fuzzers go. The fuzzers described herein may be further developed to add features such as individual probabilities for expansion rules, and depth control on a key-specific basis. This paper also assumes existing well-defined grammars for all use cases, whereas in actual application, we may instead be required to infer input grammars from sample input data corpora. These changes are likely required for the fuzzers to be considered for use in production or real testing.

### 9.3 Grammar Cheapening

The approach we have taken to cheapening grammars in this paper, while more efficient/less wasteful than alternative approaches we studied, is susceptible to creating dead-ends in grammars that contain definitions with no non-recursive rules. To minimally repopulate the cheap grammar such that no dead-ends are left we have, as an additional step, used the cheapening approach implemented in the F1 fuzzer [Gopinath and Zeller, 2019] which assigns computational cost to each rule in a grammar, and chooses rules with minimum cost of all rules in a definition.

<sup>1</sup>Generated from the ANTLR .g4 format JSON grammar (<https://raw.githubusercontent.com/antlr/grammars-v4/master/json/JSON.g4>)

## 10 Conclusion

Fuzzing is a vastly useful technique in the process of security-testing software applications. Improvements in the speed of and control over fuzzed output generation could potentially optimize this process in a significant way; this paper has demonstrated how one might contribute to both areas of improvement; optimizing fuzzers by compiling fuzzers from grammar definitions - our fuzzers reach levels of throughput comparable to some of the faster fuzzers that currently exist - and by fuzzing from grammars to generate valid inputs much faster than a random fuzzer.

## References

- [Diehl, 2014] Diehl, C. (2014). Dharma. <https://blog.mozilla.org/security/2015/06/29/dharma/>, Last accessed on 28-01-2024.
- [Felderer et al., 2016] Felderer, M., Büchler, M., Johns, M., Brucker, A. D., Breu, R., and Pretschner, A. (2016). Security testing. *Advances in Computers*, 101:1–51.
- [Gopinath and Zeller, 2019] Gopinath, R. and Zeller, A. (2019). Building fast fuzzers. *ArXiv*, abs/1911.07707.
- [Hodován et al., 2018] Hodován, R., Kiss, A., and Gyimóthy, T. (2018). Grammarinator: a grammar-based open source fuzzer. In *Proceedings of the 9th ACM SIGSOFT International Workshop on Automating TEST Case Design, Selection, and Evaluation*, A-TEST 2018, page 45201348, New York, NY, USA. Association for Computing Machinery.
- [Salem and Song, 2019] Salem, H. A. A. and Song, J. (2019). A review on grammar-based fuzzing techniques. *International Journal of Computer Science and Security (IJCSS)*, 13(3):114 – 123.

## A Source Code & Resources

The source code and grammars for the fuzzers built in this paper are contained in <https://github.com/Positron11/fuzzer>.

The G4 grammars that the JSON grammars were compiled from can be found at <https://github.com/antlr/grammars-v4/>.

The tools used to perform the aforementioned compilation were provided by <https://github.com/vrthra/antlr2json/>